

▼ Hill and Valley Prediction using Logistic Regression

Get Understanding about Data set

Each record represents 100 points on a two-dimensional graph. When plotted in order (from 1 through 100) as the Y coordinate, the points will create either a Hill (a "bump" in the terrain) or a Valley (a "dip" in the terrain). See the original source for some examples of these graphs.

1-100: Labeled "V##", Floating point values (numeric), the X-values. 101: Labeled "Class", Binary {0,1} representing {valley, hill}.

▼ Import Library

```
import pandas as pd
```

```
import numpy as np
```

▼ Import CSV as DataFrame

```
df = pd.read_csv(r'/content/Hill Valley Dataset.csv')
```

Saved successfully!

▼ Get the First Five Rows of DataFrame

```
df.head()
```

	V1	V2	V3	V4	V5	V6	V7	V8	V9
0	39.02	36.49	38.20	38.85	39.38	39.74	37.02	39.53	38.81
1	1.83	1.71	1.77	1.77	1.68	1.78	1.80	1.70	1.75
2	68177.69	66138.42	72081.88	74304.33	67549.66	69367.34	69169.41	73268.61	74465.84

▼ Get Information of DataFrame

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1212 entries, 0 to 1211
Columns: 101 entries, V1 to Class
dtypes: float64(100), int64(1)
memory usage: 956.5 KB
```

▼ Get the Summary Statistics

```
df.describe()
```

	V1	V2	V3	V4	V5	V6
count	1212.000000	1212.000000	1212.000000	1212.000000	1212.000000	1212.000000
mean	8169.091881	8144.306262	8192.653738	8176.868738	8128.297211	8173.037211
std	17974.950461	17881.049734	18087.938901	17991.903982	17846.757963	17927.117963
min	0.920000	0.900000	0.850000	0.890000	0.880000	0.860000
25%	301.425000	295.205000	297.260000	299.720000	295.115000	294.385000
75%	5358.795000	5417.847500	5393.367500	5388.482500	5321.987500	5328.047500
max	117807.870000	108896.480000	119031.350000	110212.590000	113000.470000	116848.350000

8 rows × 101 columns



▼ Get Column Names

Define y (dependent or label or target variable) and X (independent or features or attribute Variable)

```
y = df['Class']
```

```
y.shape
```

```
(1212,)
```

```
y
```

```
0      0
1      1
2      1
3      0
4      0
..
1207    1
1208    0
1209    1
1210    1
1211    0
```

```
Name: Class, Length: 1212, dtype: int64
```

```
X = df.drop('Class', axis=1)
```

```
X.shape
```

```
(1212, 1000)
```

Saved successfully!




```
X
```

	V1	V2	V3	V4	V5	V6	V7	V8	
0	39.02	36.49	38.20	38.85	39.38	39.74	37.02	39.53	38.85
1	1.83	1.71	1.77	1.77	1.68	1.78	1.80	1.70	1.77
2	68177.69	66138.42	72981.88	74304.33	67549.66	69367.34	69169.41	73268.61	74465.11
3	44889.06	39191.86	40728.46	38576.36	45876.06	47034.00	46611.43	37668.32	40980.11
4	5.70	5.40	5.28	5.38	5.27	5.61	6.00	5.38	5.28
...


▼ Get Plot of First Two Rows

```
import matplotlib.pyplot as plt
```

```
plt.plot(X.iloc[0,:])
plt.title('Valley');
```



```
plt.plot(X.iloc[1,:])
plt.title("Hill");
```



➤ Get X Variables Standardized

Standardization of datasets is a common requirement for many machinelearning estimators implemented in scikit-learn; they might behave badly if the individual features do not more or less look like standard normally distributed data: Gaussian with zero mean and unit variance.

Saved successfully!

```
from sklearn.preprocessing import StandardScaler
```

```
ss = StandardScaler()
```

```
X = ss.fit_transform(X)
```

```
X
```

```
array([[ -0.45248681, -0.45361784, -0.45100881, ..., -0.45609618,
        -0.45164274, -0.45545496],
       [ -0.45455665, -0.45556372, -0.45302369, ..., -0.45821768,
        -0.45362255, -0.45755405],
       [  3.33983504,  3.24466709,  3.58338069, ...,  3.5427869 ,
```

```

3.27907378, 3.74616847],
...,
[ 0.11084204, 0.0505953, 0.04437307, ..., 0.12533312,
 0.04456025, 0.06450317],
[-0.45272112, -0.45369729, -0.45118691, ..., -0.45648861,
 -0.45190136, -0.45569511],
[ 0.01782872, -0.02636986, 0.05196137, ..., 0.03036056,
 0.01087365, 0.03123129]])

```

X.shape

```
(1212, 100)
```

▼ Get Train Test Split

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, stratify = y, rand
```

```
X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

```
((848, 100), (364, 100), (848,), (364,))
```

▼ Get Model Train

```
from sklearn.linear_model import LogisticRegression
```

Saved successfully!



```
lr.fit(X_train, y_train)
```

```

▼ LogisticRegression
LogisticRegression()

```

▼ Get Model Prediction

```
y_pred = lr.predict(X_test)
```

```
y_pred.shape
```

```
(364,)
```

```
y_pred
```

```
array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1,
       0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0,
       0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0,
       0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0,
       0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,
       1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0,
       0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,
       0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1,
       0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0,
       0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 1,
       0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 1, 1, 0, 0, 1,
       0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 1,
       0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
       0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 1, 0, 0, 0, 1, 0, 1,
       1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 0,
       0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0,
       0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0,
       0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1])
```

▼ Get Probability of Each Predicted Class

```
lr.predict_proba(X_test)
```

Saved successfully!




```
[0.5086123, 0.49132877],
[0.63597662, 0.36402338],
[0.50894067, 0.49105933],
[0.53221095, 0.46778905],
[0.4735918 , 0.5264082 ],
[0.51131842, 0.48868158],
[0.47711039, 0.52288961],
[0.50836868, 0.49163132],
[0.51791139, 0.48208861],
[0.50929629, 0.49070371],
[0.77824423, 0.22175577],
[0.50878251, 0.49121749],
[0.53444807, 0.46555193],
[0.68401993, 0.31598007],
[0.50837605, 0.49162395],
[0.49955891, 0.50044109],
[0.66733466, 0.33266534],
[0.51287013, 0.48712987],
[0.5100684 , 0.4899316 ],
[0.49493863, 0.50506137],
[0.50415413, 0.49584587],
[0.50875424, 0.49124576],
[0.50859968, 0.49140032],
[0.51714632, 0.48285368],
[0.49828539, 0.50171461],
[0.50868242, 0.49131758],
[0.5887849 , 0.4112151 ],
[0.50578145, 0.49421855],
[0.50864404, 0.49135596],
[0.78691964, 0.21308036],
[0.50865314, 0.49134686],
[0.5085737 , 0.4914263 ],
[0.48168718, 0.51831282],
[0.50838591, 0.49161409],
[0.05333433, 0.94666567]])
```

Get Model Evaluation

Saved successfully!

```
from sklearn.metrics import confusion_matrix, classification_report
```

```
print(confusion_matrix(y_test, y_pred))
```

```
[[181  1]
 [106 76]]
```

```
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.63	0.99	0.77	182
1	0.99	0.42	0.59	182

accuracy			0.71	364
macro avg	0.81	0.71	0.68	364
weighted avg	0.81	0.71	0.68	364

▼ Get Future Predictions

Lets select a random sample from existing dataset as new value

Steps to follow

1. Extract a random row using sample function
2. Seperate X and y
3. Standardize X
4. Predict

```
X_new = df.sample(1)
```

X_new

	V1	V2	V3	V4	V5	V6	V7	V8	V9	V10	...	V92	V93	V94	V95	V96
88	1.25	1.19	1.25	1.19	1.2	1.18	1.18	1.18	1.14	1.14	...	1.19	1.17	1.08	1.04	1.05

1 rows × 101 columns



Saved successfully!



```
X_new.shape
```

```
(1, 101)
```

```
X_new = X_new.drop("Class", axis = 1)
```

X_new

	V1	V2	V3	V4	V5	V6	V7	V8	V9	V10	...	V91	V92	V93	V94	V95
88	1.25	1.19	1.25	1.19	1.2	1.18	1.18	1.18	1.14	1.14	...	1.19	1.19	1.17	1.08	1.04

4 rows x 100 columns

X_new.shape

(1, 100)

X_new = ss.fit_transform(X_new)

y_pred_new = lr.predict(X_new)

y_pred_new

array([1])

lr.predict_proba(X_new)

↗ array([[0.49714993, 0.50285007]])

Saved successfully!